**Eduardo Cabrera-Lopez**

**ITAI 1378**

**Professor Patricia McManus**

# Reflective Journal: CNN-based Classification Task

## CNN Architecture

During this lab, I had the opportunity to work with Convolutional Neural Networks (CNNs) and explore how they differ from traditional neural networks. A CNN is specifically designed to process image data by recognizing patterns such as edges and textures. The key components of a CNN are the convolutional layers, pooling layers, and fully connected layers. For example, when a convolutional layer applies filters (or kernels) to an image, it can extract features like vertical or horizontal edges. Pooling layers then down sample these feature maps to reduce the complexity while still retaining important information. In contrast, traditional neural networks do not have this ability to capture spatial relationships within an image since they treat input data as a flat vector. This is what makes CNNs more suitable for image-related tasks.

## Model Performance

I was really impressed with how my CNN model performed, reaching about 85% accuracy during training. It did a good job distinguishing between chihuahuas and muffins, though there were a few instances where it mistakenly identified muffins as chihuahuas. This mostly happened when the muffins had textures, like frosting, that resembled fur. These errors made me appreciate how challenging it can be for CNNs to handle objects with similar visual characteristics.

One thing that really caught my attention was how the model's performance improved over time as I adjusted the number of epochs. At first, with fewer epochs, the model was underfitting, meaning it wasn't learning enough from the data to detect more complex features. However, when I increased the number of epochs, the model's accuracy improved significantly. I had to be mindful of overfitting, though, as I noticed a slight dip in validation accuracy when the model became too tailored to the training data.

## Comparison

When comparing CNNs to the traditional neural network we worked with earlier, the CNN was clearly more efficient and accurate. In the previous workshop, the fully connected layers of the traditional neural network took much longer to process image data, and the accuracy was considerably lower. Without convolutional layers, the model struggled to detect the subtle patterns within the images, resulting in poor performance. In contrast, the CNN's ability to focus on specific regions of the image, extract key features, and then reduce data complexity through pooling layers led to faster training times and better overall results.

One of the major challenges I faced in this lab was deciding how many convolutional layers to include and what filter sizes to use. Initially, I started with only one convolutional layer, but the model wasn't capturing enough detail, leading to suboptimal performance. After reviewing course materials, I added a second convolutional layer and experimented with different filter sizes, which significantly improved the accuracy. Another challenge was balancing the number of epochs; at first, my model underfit, but after increasing the number of epochs and using early stopping, I was able to prevent overfitting and improve the generalization.

## Real-World Applications

This lab really opened my eyes to the wide range of practical uses for CNNs. One key area where CNNs are making a big impact is healthcare. For example, they're used in medical imaging to help detect diseases like cancer by analyzing MRI scans. CNNs are also crucial in the development of autonomous vehicles, where they handle tasks like object detection and classification, enabling cars to recognize pedestrians, other vehicles, and road signs. It's impressive how these models can be applied to such critical and diverse fields.

One application I found particularly fascinating is how CNNs are transforming the retail industry by improving image search algorithms. For example, many online shopping platforms now allow customers to search for products simply by uploading a picture rather than typing out a description. This is incredibly useful when shoppers don't know exactly how to describe what they're looking for but have a visual reference. CNNs make this possible by accurately recognizing and classifying images, matching them to relevant products. Beyond retail, this kind of image recognition technology is proving to be essential in a wide range of fields, from fashion and interior design to even inventory management, where visual data can streamline operations. It's exciting to see how CNNs are making everyday tasks more efficient and accessible.

## Ethical Factors

While CNNs offer powerful tools for solving real-world problems, there are ethical concerns to consider. One major issue is bias in the training data. If the dataset used to train the model is unbalanced, the model may develop biases, leading to unfair or inaccurate results. This is particularly concerning in fields like facial recognition, where biases in race or gender can have significant consequences.

Privacy is another ethical consideration. Models that rely on personal data, such as those used in facial recognition or surveillance, must ensure that individuals' privacy is protected. With regulations like GDPR, it is important to consider how data is collected and used to avoid potential misuse.

In conclusion, this lab on CNN-based image classification provided me with a deeper understanding of how convolutional layers and filters work together to detect and classify patterns in images. Through hands-on experimentation, I gained insights into the strengths of CNNs, particularly in their ability to process complex visual data efficiently. The challenges I encountered, such as tuning the model's architecture and preventing overfitting, helped reinforce my learning and problem-solving skills. Overall, this experience has solidified my appreciation for the power of CNNs in real-world applications, from healthcare to autonomous driving, while also reminding me of the ethical considerations that must be addressed when developing AI models.

**References:**


Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep learning*. MIT Press.

Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet classification with deep convolutional neural networks. *Communications of the ACM, 60*(6), 84-90. https://doi.org/10.1145/3065386

LeCun, Y., Bengio, Y., & Hinton, G. (2015). Deep learning. *Nature, 521*(7553), 436–444. https://doi.org/10.1038/nature14539

Rawat, W., & Wang, Z. (2017). Deep convolutional neural networks for image classification: A comprehensive review. *Neural Computation, 29*(9), 2352-2449. https://doi.org/10.1162/neco_a_00990